

PRODUCT

May 20, 2026

Governing AI agents at scale with Unity Catalog

Four pillars for governing every model call, tool invocation, and agent interaction in your organization

by [David Nasi](#) and [Stefania Leone](#)

SUMMARY



A year ago, your organization had a dozen AI agents. Today, there are thousands.

Every developer has a coding agent that writes, reviews, and ships code alongside them. Your analytics team built forecasting agents. Sales operations deployed lead scoring. The Support organization automated ticket routing. Marketing launched personalization. Finance built reconciliation workflows. Every team saw an opportunity and moved fast.

Now someone asks: "Which agents are accessing customer PII?"

The answer requires pulling logs from dozens of systems, manually correlating them, and hoping nothing was missed. Each agent logs, authenticates and accesses data differently. There's no single place to look.

Or maybe you took the opposite path. You locked everything down. No agents got deployed without extensive review. Security stayed tight. But now you're six months behind competitors who moved faster. Developers and users are frustrated. Some have left for companies where they can actually use AI tools.

Neither extreme works. Ungoverned agents create risk you can't measure. Locked-down environments create a different kind of risk: falling behind while talent walks